

2026 edition

Deep Learning for Music Analysis and Generation

Introduction

Intros + course logistics



Yi-Hsuan Yang Ph.D.

yhyangtw@ntu.edu.tw

Course Website

<https://cool.ntu.edu.tw/courses/65378>

<https://affige.github.io/teaching.html>

<https://github.com/affige/DeepMIR/>




Outline

- **Intros of Music & AI**
- Course logistics
- Fundamentals of musical audio
- Appendix: ML 101


Music AI; or *Music Information Research* (MIR)

- Music analysis



- music understanding
- music search
- music recommendation


- Music generation



- MIDI generation
- audio generation
- MIDI-to-audio generation


Music AI: {signal processing, machine learning} + music

- Music analysis



- music understanding
- music search
- music recommendation

- Music generation



- MIDI generation
- audio generation
- MIDI-to-audio generation

Music Analysis Demo 1: Music Transcription

<https://songscription.ai/>



- songscription.ai
- AnthemScore
- Melody Scanner
- Klangio

Flight of the Bumblebee
Nikolai Rimsky-Korsakov

Up Down Scale
Exercise

Invention No. 4 in D Minor
Johann Sebastian Bach

Flight of the Bumblebee
Nikolai Rimsky-Korsakov

Original Audio

0:00 / 0:22

Play Generated Sheet Music

Flight-of-the-bumblebee-piano


A musical score for the piece 'Flight of the Bumblebee' by Nikolai Rimsky-Korsakov. The score is displayed in a piano arrangement, featuring a treble and bass staff. The melody is written in the treble staff, and the bass staff provides harmonic support. The score is titled 'Flight-of-the-bumblebee-piano' and includes a play button and a progress bar at the bottom.

Music Analysis Demo 2: Source Separation

<https://www.youtube.com/watch?v=r4NcZbgtlgw>



- Apple Logic Pro
- iZotope RX
- Gaudio Studio
- RipX
- moises.ai
- lalal.ai



Source: <https://support.apple.com/en-md/guide/logicpro/lgcp61bae908/mac>

Music Analysis Demo 3: Audio Understanding

<https://research.nvidia.com/labs/adlr/AF3/>



Audio Flamingo 3


Audio	Instruction	Ground-truth	Audio Flamingo 3
	What is surprising about the relationship between the barking and the music?	Dogs barking in sync with the music	The barking and music are synchronized , creating a rhythmic and harmonious interaction.
	What makes this saxophone performance unlike typical saxophone sounds?	Aside from "Interstellar Spaces", Saxophones do not bark nor meow	The saxophone is played in a way that resembles a dog howling , creating a unique and unusual sound.
	What makes this combination of acoustic and ambient sounds unlikely to appear together?	Unlikely existence of banjo and rain sounds at the same time.	The contrast between the lively music and the sound of raindrops creates an unusual auditory experience.

Positive Use Cases of Music Generation

- **Make music easier to play with** (i.e., for common people)
 - “democratization” of music creation
- **Make musicians’ life easier** (i.e., for musicians)
 - inspire ideas
 - suggest continuations
 - suggest accompaniments
- **Create copyright free music for videos or games**
- **Music education / learning**
- **Digital archival**

Evolution of Music Generation Research

- Monophonic MIDI (melody) → polyphonic MIDI (e.g. piano)
- Single-instrument MIDI → multi-instrument MIDI
- MIDI → audio
- Solo → full song




Source: <https://www.mdpi.com/2079-9292/14/6/1197>

Music Generation Demo 1: Piano Genie

(Make music easier to play with)

<https://magenta.tensorflow.org/pianogenie>



Music Generation Demo 2: Tone Transfer

(Make music easier to play with)

<https://sites.research.google/tonetransfer>




Music Generation Demo 3: Song Generation

(Create copyright free music; for musicians)

<https://www.youtube.com/watch?v=MCLVhRjlvys>




moises




Music Generation Demo 4: Song Generation

(Create copyright free music; for common people)

<https://suno.com/song/722aa074-5010-44d1-ad51-55468f22abc2>



<https://suno.com/>



<https://www.mureka.ai/>

Global Interest in Music AI: Industry



(Slide from Rujing Huang, Bob L. T. Sturm, and Andre Holzapfel, "De-centering the West: East Asian Philosophies and the Ethics of Applying Artificial Intelligence to Music," ISMIR 2021)

Global Interest in Music AI: Academia

- MIT EECS | Anna Huang (<https://czhuang.github.io/>)
- CMU CS | Chris Donahue (<https://chrisdonahue.com/>)
- U Michigan SMTD | Hao-Wen Dong (<https://hermandong.com/>)
- Stanford CCRMA | Ge Wang (<https://ccrma.stanford.edu/~ge/>)
- NYU MARL | Brian McFee (<https://brianmcfee.net/>)
- U Rochester | Zhiyao Duan (<https://hajim.rochester.edu/ece/sites/zduan/>)
- Georgia Tech | Alexander Lerch (<https://www.alexanderlerch.com/>)
- Centre for Digital Music, Queen Mary University of London
(<https://www.c4dm.eecs.qmul.ac.uk/>)
- Academia Sinica | Li Su (<https://homepage.iis.sinica.edu.tw/pages/lisu/>)

Outline

- Intros of Music & AI
- **Course logistics**
- Fundamentals of musical audio
- Appendix: ML 101

This Course: Prerequisites


- **Graduate level (CommE5070) @ NTU GICE**
 - It's **NOT** a music course
 - It's an EE/CS graduate-level course working on music data/problems
- Suitable for people who have
 - Great interest in music
 - Good background in machine learning & deep learning
 - Good coding experience in python and a deep learning framework such as PyTorch
- Not a class designed for deep learning beginners

This Course: Wills and Won'ts

- Will talk about
 - Domain knowledge in music data representation
 - Domain knowledge in music analysis: timbre, rhythm, pitch
 - Deep learning-based music analysis
 - Deep learning-based audio generation
 - Deep learning-based MIDI sequence generation
- **Won't** talk about (too much on)
 - Basics in machine learning and deep learning
 - Applications in other domains

This Course: Objective

- Get you ready to do **ML** research on music AI
- Lots of hands-on



Lecturer

- Lecturer
 - Yi-Hsuan Yang (楊奕軒)
 - <https://affige.github.io/>
 - yhyangtw@ntu.edu.tw
- Office hour
 - Tuesday 9:30-11:30, or by appointment
 - Office: BL-616 (電二)

Teaching Assistants

- TAs
 - Ting-Yi Hu (胡庭翊)
 - tingyihu521@gmail.com
 - Bo-Rui Chen (陳柏睿)
 - rakechen168@gmail.com
- Office hour
 - Monday 13:30-15:00, or by appointment
 - Office: BL-505 (博理館)

Location & Time

- Location: 電二229
- Time: Tuesday 6,7,8
- Usually
 - 6*: 13:40-14:30 (i.e., 20 mins later)
 - 7*: 14:40-15:30
 - 8*: 15:40-16:30
- Last week (W16)
 - 6: 13:20- (we will begin the class on time for the last week)

Online Option

- Live through Google Meet: <http://meet.google.com/dka-enxo-toy>

Textbook

(for the music analysis part)

- Reference textbook

Meinard Müller


Fundamentals of Music Processing
Using Python and Jupyter Notebooks

ISBN: 978-3-030-69808-9

Springer, April 2021

<https://www.audiolabs-erlangen.de/fau/professor/mueller/bookFMP>

<https://www.audiolabs-erlangen.de/resources/MIR/FMP/C0/C0.html>



- Related book




- *An Introduction to Audio Content Analysis: Applications in Signal Processing and Music Informatics*, Wiley




<https://github.com/alexanderlerch/pyACA>

<https://github.com/alexanderlerch/ACA-Slides>

FMP Notebook

<https://www.audiolabs-erlangen.de/resources/MIR/FMP/C0/C0.html>


Part	Title	Notions, Techniques & Algorithms	HTML	IPYNB
	Basics	Basic information on Python, Jupyter notebooks, Anaconda package management system, Python environments, visualizations, and other topics	[html]	[ipynb]
	Overview	Overview of the notebooks (https://www.audiolabs-erlangen.de/FMP)	[html]	[ipynb]
	Music Representations	Music notation, MIDI, audio signal, waveform, pitch, loudness, timbre	[html]	[ipynb]
	Fourier Analysis of Signals	Discrete/analog signal, sinusoid, exponential, Fourier transform, Fourier representation, DFT, FFT, STFT	[html]	[ipynb]
	Music Synchronization	Chroma feature, dynamic programming, dynamic time warping (DTW), alignment, user interface	[html]	[ipynb]

Part	Title	Notions, Techniques & Algorithms	HTML	IPYNB
	Music Structure Analysis	Similarity matrix, repetition, thumbnail, homogeneity, novelty, evaluation, precision, recall, F-measure, visualization, scape plot	[html]	[ipynb]
	Chord Recognition	Harmony, music theory, chords, scales, templates, hidden Markov model (HMM), evaluation	[html]	[ipynb]
	Tempo and Beat Tracking	Onset, novelty, tempo, tempogram, beat, periodicity, Fourier analysis, autocorrelation	[html]	[ipynb]
	Content-Based Audio Retrieval	Identification, fingerprint, indexing, inverted list, matching, version, cover song	[html]	[ipynb]
	Musically Informed Audio Decomposition	Harmonic/percussive separation, signal reconstruction, instantaneous frequency, fundamental frequency (F0), trajectory, nonnegative matrix factorization (NMF)	[html]	[ipynb]

Grading Policy

- Grading policy
 - **Assignments** (60%), 3 times, each 20%
 - **Final Project** (40%): for teams of 2 or **3** (recommended)
- Work reasonably hard to get a high score

Course Material



Topics

- **audio** → **audio**: signal processing
- **audio** → **score**: transcription
- **score** → **score**: composition
- **score** → **audio**: synthesis
- **audio** → **knowledge**: audio analysis
- **score** → **knowledge**: symbolic-domain analysis

Syllabus

- **W1.** Intros & fundamentals of musical audio
- **W2.** Classification & foundation models
 - **HW1:** Music classification
- **W3.** Large audio language models & music transcription
- **W4.** Source separation & vocoders
- **W5.** Symbolic music generation I
 - **HW2:** MIDI generation
- **W6.** Symbolic music generation II
- **W7.** Musical Audio Generation I
- **W8.** Musical Audio Generation II
 - **HW3:** Text-to-music generation
- **W9.** Musical Audio Generation III
- **W10.** Musical Audio Generation IV
- **W11.** Effect modeling & automatic mixing
- **W12. Project pitch** (proposals of final projects)
- **W13.** Miscellaneous Topics
- **W14.** Miscellaneous Topics
- (W15. Break)
- **W16. Oral presentation** of final projects

Assignments

- **Programming (in python) + report (in English)**
 - We assume that you are comfortable with programming in python and deep learning frameworks such as PyTorch
 - One new assignment **every three weeks** (W2, W5, W8); each due in 3 weeks (W5, W8, W11)
 - Submit **code + model + report**
 - Will select around 3 people to share their work
 - NO cheating: Will run *plagiarism detector*
- **Topics**
 - HW1: music classification
 - HW2: MIDI generation
 - HW3: text-to-music audio generation

Assignments

		HW1	HW2	HW3	Final project
		Classification	MIDI generation	Text-to-music	
W1	Intros + course logistics & fundamentals of musical audio				
W2	Music classification & music foundation models	announced			
W3	Large audio language models & music transcription	due			
W4	Source separation & vocoders				
W5	Symbolic music generation I		announced		
W6	Symbolic music generation II		due		
W7	Musical audio generation I: codec models				
W8	Musical audio generation II: latent models			announced	
W9	Musical audio generation III: lyrics-to-song generation			due	
W10	Musical audio generation IV: cover-gen & MIDI-to-audio				team up
W11	Effect modeling & automatic mixing				
W12	Project pitch				proposal
W13	Miscellaneous topics				
W14	Miscellaneous topics				
W15	break				
W16	Final presentation				oral report
W17					
W18					paper report

NO Cheating


- Once caught: **failure** of the course

Final Project

- For teams of 2 or **3** (recommended)
- Start earlier & form teams
- Deadline for **team-up**: W10
- **Project pitch**: W12
- **Final presentation**: W16
- Deadline for **final report**: W16+2


Final Project

- Music analysis



- music semantic labeling
- music transcription
- source separation

- Music generation



- MIDI generation
- audio generation
- MIDI-to-audio generation


Final Project Showcase

https://affige.github.io/teaching_deepmir25.html

- **List of final projects** presented by the class students
 - *BLT: An IPA-Aware Agentic Framework for Controllable Lyric Translation*
 - extensions of this project has been accepted for publication at [EMNLP'26](#): "IPA-Verified Singable Lyrics Translation with LLM Agents"
 - *High-level Semantic Direction Discovery*
 - extensions of this project has been accepted for publication at [KDD'26](#): "AnchorSteer: Self-discovered concept injection for structure-preserving music editing"
 - *Style Transfer of Audio Effects with Differentiable Signal Processing*
 - *Unified Audio and Text-Controlled Automatic Music Mixing via CLAP Embeddings*
 - extensions of this project still ongoing
 - *Make Spoken Language Models Sing: A Preliminary Study using Text-Aligned Tokenization*
 - extensions of this project still ongoing

Final Project Showcase

- Multimodal music emotion recognition: audio + lyrics + MIDI




Song	Model	Prediction
Bruno Mars, Anderson .Paak, Silk Sonic - 777 (Happy)	CRNN	Happy
	MERT	Happy
	Whisper + BERT	Angry
	Late fusion	Angry
	LSTM-Attn + Remi	Angry

Final Project Showcase

- Guitar effect removal

https://y10ab1.github.io/guitar_effect_removal/



Final Project Showcase

- Transfer music to failed recorder style
<https://navi0105.github.io/demo/music2fail/>

genmusic_demo_list

https://github.com/affige/genmusic_demo_list

About

a list of demo websites for automatic
music generation research

artificial-intelligence

music-generation

Resources

- ML/DL
 - <http://speech.ee.ntu.edu.tw/~tlkagk/courses.html>
 - <https://www.csie.ntu.edu.tw/~htlin/course/>
 - <https://www.csie.ntu.edu.tw/~yvchen/teaching>
 - <https://courses.cs.washington.edu/courses/cse599i/20au/> (generative models)
- Music information research
 - https://www.audiolabs-erlangen.de/fau/professor/mueller/teaching/2023w_mpa
 - <https://musicinformationretrieval.com/>
 - <https://mac.kaist.ac.kr/~juhan/gct634/index.html>
 - <http://www.jordipons.me/apps/teaching-materials/>
 - <https://www.upf.edu/web/smc/audio-signal-processing-for-music-applications>

Resources

- Conference proceedings
 - **Int'l Soc. Music Information Retrieval Conf. (ISMIR)**
 - Int'l Conf. Acoustic, Speech, and Signal Processing (ICASSP)
 - ACM MM, ACM ICMR, ACM SIGIR, IEEE ICME
- Transactions
 - IEEE Trans. Audio, Speech and Language Processing (TASLP)
 - IEEE Trans. Multimedia (TMM)
 - IEEE Trans. Signal Processing (TSP)

Additional Enrollment

<https://forms.gle/8CMCrwmbh9GU69ph6>

- **Sign up** at NTU Cool (選課意願登記), if you are an NTU student
- **AND, Fill the form** before 23:59, **September 11** (Friday)
 - DL background
 - Music background
 - Ideas for final project
- Will announce the result before next **Monday**
 - Not first-come-first-serve; will read your submission
 - Will only send a mail to those selected




Outline

- Intros of Music & AI
- Course logistics
- Fundamentals of musical audio
- **Appendix: ML 101**


Machine Learning

<https://www.microsoft.com/en-us/research/people/cmbishop/prml-book/>




Freely available

1. Introduction
2. Probability Distributions
3. Linear Models for Regression
4. Linear Models for Classification
5. Neural Networks
6. ...





Classification: Overview (1/2)



- Given: $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$
 - $\mathbf{x}_i \in \mathbb{R}^M$: feature representation of a data instance
 - $y_i \in \{-1, +1\}$: class label of that instance
- Goal
 - learn a separating hyperplane (decision boundary)

Classification: Overview (2/2)

- Given: $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$
 - $\mathbf{x}_i \in \mathbb{R}^M$: **feature** representation
 - $y_i \in \{-1, +1\}$: class **label**



Model Training (1/3)

1. Collect data (y) and extract features ($\mathbf{x}$)
2. Build model: choose hypothesis class $\mathcal{H}$ and loss function l
 - hypothesis class
 - **linear:** $\hat{y} = f(\mathbf{x}, \mathbf{w}, b) = \text{sgn}(\mathbf{w}^T \mathbf{x} + b)$
 - **quadratic:** $\hat{y} = f(\mathbf{x}, \mathbf{A}, \mathbf{w}, b) = \text{sgn}(\mathbf{x}^T \mathbf{A} \mathbf{x} + \mathbf{w}^T \mathbf{x} + b)$

$$\text{sgn}(x) = \begin{cases} +1, & x \geq 0; \\ -1, & x < 0. \end{cases}$$

□ “parameters” to be *learned*: $\theta = \{\mathbf{A}, \mathbf{w}, b\}$

Model Training (2/3)

1. Collect data (y) and extract features (x)
2. Build model: choose hypothesis class $\mathcal{H}$ and loss function l

- loss function

- **indicator function:** $l(y, \hat{y}) = 0$ if $y = \hat{y}$;
 $l(y, \hat{y}) = 1$ otherwise

- **square loss:** $l(y, \hat{y}) = (1 - y\hat{y})^2$

- **hinge loss:** $l(y, \hat{y}) = \max(0, 1 - y\hat{y})$

- ◻ while $y \in \{-1, +1\}$, $\hat{y}$ may not

- ◻ hinge loss is used in support vector machine (SVM)

Model Training (3/3)


1. Collect data (y) and extract features (x)
2. Build model: choose hypothesis class $\mathcal{H}$ and loss function l
3. **Optimization**: minimize the empirical loss
 - empirical loss: $\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N l(y_i, f(x_i, \theta))$
 - optimization: $\min_{\theta} \mathcal{L}(\theta) \Rightarrow$ by computing $\frac{\partial \mathcal{L}(\theta)}{\partial \theta}$

More on Loss Functions (1/4)

- Different loss function leads to different machine learning problems
 - **classification**: $l(y, \hat{y}) = (1 - y\hat{y})^2$
 - **regression**: $l(y, \hat{y}) = (y - \hat{y})^2$

Given N inputs $(\mathbf{x}_i, y_i)$, $1 \leq i \leq N$, where $\mathbf{x}_i$ is feature and y_i is the numerical value to be predicted, train a regression model $f(\cdot)$ such that the **mean squared error (MSE)** is minimized


$$\min_f \frac{1}{N} \sum_{i=1}^N (y_i - f(\mathbf{x}_i))^2$$



More on Loss Functions (2/4)

- Different loss function leads to different machine learning problems
 - **classification**: $l(y, \hat{y}) = (1 - y\hat{y})^2$
 - **regression**: $l(y, \hat{y}) = (y - \hat{y})^2$
 - **clustering** (an unsupervised task)


$$\mathcal{L}(\Delta) = \sum_{k=1}^K \sum_{i \in C_k} \|x_i - \mu_k\|^2$$



More on Loss Functions (3/4)

- Adding regularizers

- original: $\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{w}^T \mathbf{x}_i)^2$
- modified: $\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \|\mathbf{w}\|_2^2$



(b) ℓ_2 -ball meets quadratic function. ℓ_2 -ball has no corner. It is very unlikely that the meet-point is on any of axes.

$$\|\mathbf{w}\|_2^2 = \mathbf{w}^T \mathbf{w} = \sum_{j=1}^M w_j^2$$


prefer $\mathbf{w} = [1.5, -0.5, 0.8, -1.0]^T$
over $\mathbf{w} = [9.5, -8.0, 2.5, -3.0]^T$,
despite that the latter has smaller MSE

- $\lambda \geq 0$ is a parameter to be tuned (i.e. empirically determined)

More on Loss Functions (4/4)

- Adding regularizers

- modified: $\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \|\mathbf{w}\|_2^2$
- modified: $\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \|\mathbf{w}\|_1$



(a) ℓ_1 -ball meets quadratic function. ℓ_1 -ball has corners. It's very likely that the meet-point is at one of the corners.

$$\|\mathbf{w}\|_1 = \sum_j |w_j|$$

prefer $\mathbf{w} = [1.5, 0, 0, -2.0]^T$
over $\mathbf{w} = [1.5, -0.5, 0.8, -1.0]^T$,
despite that the latter has smaller MSE

- L1 is the only norm that is both convex and sparsity promoting

Example: Ridge Regression (for Regression)

- Loss function

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{w}^T \mathbf{x}_i)^2 + \lambda \|\mathbf{w}\|_2^2$$

$$\text{or, } \mathcal{L}(\mathbf{w}) = (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w}$$

$\begin{matrix} \swarrow & \swarrow & \searrow \\ N \times 1 & N \times M & M \times 1 \end{matrix}$

- Optimization

$$\frac{\partial \mathcal{L}(\mathbf{w})}{\partial \mathbf{w}} = \underset{M \times N}{2\mathbf{X}^T} \underset{N \times 1}{(\mathbf{X}\mathbf{w} - \mathbf{y})} + 2\lambda \underset{M \times 1}{\mathbf{w}} \stackrel{\text{let}}{=} \mathbf{0}$$

$$\therefore \mathbf{X}^T \mathbf{y} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}) \mathbf{w}$$

$$\therefore \mathbf{w} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$


Example: Support Vector Machine (1/3)

- Loss function

$$\min_{\mathbf{w}, \xi_i} \|\mathbf{w}\|_2 + C \sum_{i=1}^N \xi_i$$


subject to $y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i$ for $i = 1 \dots N$

- the idea of hinge loss $\max(0, 1 - y\hat{y})$ is also to have $y\hat{y} \geq 1$
- the slack variables $\xi_i \geq 0$ allow for errors



Example: Support Vector Machine (2/3)

- Loss function: $\min \|\mathbf{w}\|_2 + C \sum_{i=1}^N \xi_i$
 - smaller C : reduced model complexity
 - larger C : less misclassification
- Demo: <https://www.csie.ntu.edu.tw/~cjlin/libsvm/>




Example: Support Vector Machine (3/3)

- Different kernel functions of SVM
 - the **RBF** kernel (-t 2) is often used for nonlinear SVM; important parameters: C and γ


```
options:
-s svm_type : set type of SVM (default 0)
    0 -- C-SVC
    1 -- nu-SVC
    2 -- one-class SVM
    3 -- epsilon-SVR
    4 -- nu-SVR
-t kernel_type : set type of kernel function (default 2)
    0 -- linear:  $u \cdot v$ 
    1 -- polynomial:  $(\gamma u \cdot v + \text{coef0})^{\text{degree}}$ 
    2 -- radial basis function:  $\exp(-\gamma |u-v|^2)$ 
    3 -- sigmoid:  $\tanh(\gamma u \cdot v + \text{coef0})$ 
-d degree : set degree in kernel function (default 3)
-g gamma : set gamma in kernel function (default  $1/\text{num\_features}$ )
-r coef0 : set coef0 in kernel function (default 0)
-c cost : set the parameter C of C-SVC, epsilon-SVR, and nu-SVR (default 1)
```

Example: Random Forest







- Ensemble of decision trees
 - important parameters: number of trees, tree depth
 - strength: interpretability, non-linear





Depth-2 decision tree



binary classification





multi-class classification



Example: k-Nearest Neighbors

- Classify an instance by a **majority vote** of the class membership of its k nearest neighbors
 - k is a positive integer
 - performance depends on the value of k and the features
 - lazy learning
 - nonlinear classifier



Evaluation (1/3)

- Model training
 1. Collect data and extract features
 2. Build model: choose hypothesis class and loss function
 3. Optimization: minimize the empirical loss
- **Evaluation**
 4. Evaluate on unseen, testing data

Evaluation (2/3)

- Performance measures
 - classification accuracy (CA)
 - per-class precision (P), recall (R), F-score (F)
 - confusion table

		Predicted		
		Cat	Dog	Rabbit
Actual class	Cat	5	3	0
	Dog	2	3	1
	Rabbit	0	2	11

$$CA = (5+3+11)/(5+3+0+2+3+1+0+2+11)$$

$$P(cat) = 5/(5+2)$$

$$R(cat) = 5/(5+3)$$

Evaluation (3/3)


- Performance measures
 - precision (**P**): $TP/(TP+FP)$
 - recall (**R**): $TP/(TP+FN)$
 - F-score (**F**): $2PR/(P+R)$

How many selected items are relevant?

$$\text{Precision} = \frac{\text{green semi-circle}}{\text{green and red semi-circle}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{green semi-circle}}{\text{green rectangle}}$$



Train, Validation, Test (1/2)

- Test set is unseen, how to evaluate our classifier?
 - use part of the training set as the “**validation set**”
 - e.g. 80% for training, 20% for validation
 - use the validation set to: 1) tune the parameters, 2) select classification algorithms, and 3) select features
- **K-fold Cross validation**
 - divide the data into K folds, use one fold for validation and the rest of training, repeat K times and report the average result

Train, Validation, Test (2/2)

- **Model overfitting**
 - training accuracy $\neq$ validation accuracy
 - that's why we need a balance between model complexity and training error
- **Cross-dataset generalizability**
 - validation accuracy $\neq$ test accuracy
- **Cheating issue**
 - make sure there is **no overlaps** between the training and test sets
 - **song-level split** (preferred) vs. chunk-level split

Data Labeling (1/2)

- Collect data and extract features
- How to get the labels?
 - **existing datasets**
 - due to copyright restrictions, it's not easy to get the audio files
 - **crawl the web**
 - unsure of the quality; do not know how the labeling was done
 - **manually annotate**
 - labor intensive and time consuming but can be worthwhile

Data Labeling (2/2)

- Issues to be taken care of while building a new dataset
 - 1. be aware of “confounds”**
 - Example 1: Jazz from the 60’s vs. contemporary Rock
 - Example 2: Happy songs from artist 1 vs. sad songs from artist 2
 - 2. try to make the annotation process easier**
 - binary decision -> multiple choices -> rating
 - gamification
 - 3. clear instructions**
 - 4. check for reliability/consistency**

Feature Extraction and Processing

- Collect data and extract features
- The usual pipeline
 - frame-level **feature extraction**
 - **pooling** (from frame-level to clip-level)
 - **feature normalization**
- Classification tasks in MIR
 - **clip-level prediction**: genre classification
 - **frame-level prediction**: onset detection, pitch detection
 - for clip-level prediction, pooling is often used

Feature Pooling

- Rationale
 - thousands of frames per clip → thousands of feature vectors per clip
 - we can either train a frame-level classifier, but eventually still have to integrate the result to the clip-level
 - or, we can integrate the feature vectors before training the classifier
- Approach
 - take statistics such as **mean** and **standard deviation**

Feature Normalization/Scaling

- For each individual feature dimension
 - **linear normalization:** $x' = \frac{x - \min(x)}{\max(x) - \min(x)}$
 - **z-score normalization:** $x' = \frac{x - \text{mean}(x)}{\text{std}(x)}$
 - the values $\text{mean}(x)$ and $\text{std}(x)$ are computed from all the *training* instances (i.e. *not* including the validation and the test sets)
- For each individual instance
 - **scaling to unit length:** $x' = \frac{x}{|x|}$

Early Fusion or Late Fusion

- You can extract features of different types to characterize different aspects of the signal
 - timbre, pitch, rhythm, lyrics, etc
- **Early fusion** (feature-level fusion)
 - concatenate them into a single feature vector then train a classifier
- **Late fusion** (decision-level fusion)
 - train a classifier for each feature type, and then aggregate the estimate of these classifiers (e.g., majority vote or average) to make the final decision

Feature Selection

- Not all features would be relevant
- Feature selection according to
 - domain knowledge
 - empirical observation at the feature matrix (e.g., remove a feature if there is NaN or outlier [or replace the NaNs by the average])
 - statistical methods

Library: Torchaudio

https://pytorch.org/audio/0.11.0/tutorials/audio_feature_extractions_tutorial.html

```

    waveform, sample_rate = get_speech_sample()

    n_fft = 1024
    win_length = None
    hop_length = 512

    # define transformation
    spectrogram = T.Spectrogram(
        n_fft=n_fft,
        win_length=win_length,
        hop_length=hop_length,
        center=True,
        pad_mode="reflect",
        power=2.0,
    )
    # Perform transformation
    spec = spectrogram(waveform)

    print_stats(spec)
    plot_spectrogram(spec[0], title="torchaudio")

```

Python: Librosa

<https://librosa.org/doc/main/feature.html>

Feature extraction

- ⊕ Spectral features
- ⊕ Rhythm features
- ⊕ Feature manipulation
- ⊕ Feature inversion

Onset detection

Beat and tempo

Spectrogram decomposition

Effects

Temporal segmentation

Sequential modeling

Utilities

Python: scikit-learn

<https://scikit-learn.org/1.4/tutorial/index.html>

[Install](#) [User Guide](#) [API](#) [Examples](#) [Community](#) [More ▾](#)

[Prev](#) [Up](#) [Next](#)

scikit-learn 1.4.2
[Other versions](#)

Please [cite us](#) if you use the software.

Welcome to scikit-learn
scikit-learn Tutorials
An introduction to machine learning with scikit-learn
A tutorial on statistical-learning for scientific data processing
Working With Text Data
Choosing the right estimator
External Resources, Videos and Talks
Getting Started
User Guide
Glossary of Common Terms and API Elements
Examples

This is documentation for an old release of Scikit-learn (version 1.4). Try the [latest stable release](#) (version 1.5) or [development](#) (unstable) versions.

scikit-learn Tutorials

An introduction to machine learning with scikit-learn

- Machine learning: the problem setting
- Loading an example dataset
- Learning and predicting
- Conventions

Practical Issues 1: Multi-label Classification

- A data instance can be labeled with more than one classes
- *Multi-class* classification (one-hot; one out of many) vs. *multi-label* classification (multi-hot)
- Harder for ML (may need multiple binary classifiers) but simple for DL: it's just a matter of changing the loss function
 - Multi-class classification: **categorical cross entropy** (CE)
 - Multi-label classification: **binary cross entropy** (BCE)

Practical Issues 2: Data Imbalance

- Simply predicting everything as the majority class would already give you good accuracy
- A naïve workaround: **subsampling the majority class**

Practical ML Hints (not for DL)

- Look into the features
- Remember to normalize the features
- Use validation set to properly tune the parameters
- Nonlinear classifier usually performs better, but also requires longer training time
- Use a subset of data to fast prototype your ideas first, before using the full set for more complete experiments
- Study the classification result (e.g., confusion table)